

## Tidy data

Peter Ralph

13 October – Advanced Biological Statistics

1

## Tidy data

2 . 1

### Rules of thumb for data tidiness

- Store a copy of data in a nonproprietary format, such as plain ASCII text
- Leave an uncorrected file when doing analyses
- Use descriptive names for your data files and variables
- Include a header line with descriptive variable names
- Maintain effective metadata about the data (a README)
- Add new observations to a dataset by *row*
- Add new variables to a dataset by *column*
- A column of data should contain only one data type
- All measurements of the same type should be in the same column

2 . 2

images of lab notebooks pasted into an Excel document

2 . 3

Number of eggs laid by some chickens

breed	1	2	3	4	5	6
rhode island red	5	6	NA	NA	NA	NA
white leghorn	7	5	6	8	NA	NA
barred rock	3	2	4	4	3	4
jersey giant	5	2	8	NA	NA	NA
australorp	4	NA	NA	NA	NA	NA


2 . 4

	breed	num_eggs
11	rhode island red	5
21	rhode island red	6
12	white leghorn	7
22	white leghorn	5
32	white leghorn	6
42	white leghorn	8
13	barred rock	3
23	barred rock	2
33	barred rock	4
43	barred rock	4
53	barred rock	3
63	barred rock	4
14	jersey giant	5
24	jersey giant	2
34	jersey giant	8
15	australorp	4

2 . 5

## Exercise

Design a tidy data format for the stickleback experiment:



2 . 6

Tools for tidy data

Tidying data is *hard*!

2 . 7

Tools for tidy data

Tidying data is *hard*!

... and often requires expert input.

2 . 7

Tools for tidy data

Tidying data is *hard*!

... and often requires expert input.

Many common *data wrangling* operations are made easier by the [tidyverse](#).

2 / 7

The “tidyverse”

- packages that do many of the same things as base functions in R
- designed to do them more “cleanly”
- also includes ggplot (for “Grammar of Graphics”)

2 / 8


A tibble is a data frame

```
#> # A tibble: 234 × 11
#>   manufacturer model displ year cyl  trans drv  cty hwy fl
#>   <chr> <chr> <dbl> <int> <int> <chr> <chr> <int> <int> <chr>
#> 1   audi  a4    1.8  1999   4 auto(l5) f  18  29  p
#> 2   audi  a4    1.8  1999   4 manual(m5) f  21  29  p
#> 3   audi  a4    2.0  2008   4 manual(m6) f  20  31  p
#> 4   audi  a4    2.0  2008   4 auto(av) f  21  30  p
#> 5   audi  a4    2.8  1999   6 auto(l5) f  16  26  p
#> 6   audi  a4    2.8  1999   6 manual(m5) f  18  26  p
#> # ... with 228 more rows, and 1 more variables: class <chr>
```

**manufacturer**  
**model** - model name  
**displ** - engine displacement, in litres  
**year** - year of manufacture  
**cyl** - number of cylinders  
**Trans** - type of transmission  
**drv** - f = front-wheel drive, r = rear wheel drive, 4 = 4wd  
**cty** - city miles per gallon  
**hwy** - highway miles per gallon  
**fl** - fuel type  
**class** - “type” of car

## A tibble is a data frame

```
#> # A tibble: 234 × 11
#>   manufacturer model displ year   cyl trans  drv  cty   hwy fl
#>   <chr>    <chr> <dbl> <int> <int>   <chr> <chr> <int> <int> <chr>
#> 1   audi    a4     1.8  1999     4 auto(l5) f  18  29   p
#> 2   audi    a4     1.8  1999     4 manual(m5) f  21  29   p
#> 3   audi    a4     2.0  2008     4 manual(m6) f  20  31   p
#> 4   audi    a4     2.0  2008     4 auto(av) f  21  30   p
#> 5   audi    a4     2.8  1999     6 auto(l5) f  16  26   p
#> 6   audi    a4     2.8  1999     6 manual(m5) f  18  26   p
#> # ... with 228 more rows, and 1 more variables: class <chr>
```



## Key functions in dplyr

- Pick observations by their values with `filter()`.
- Reorder the rows with `arrange()`.
- Pick variables by their names with `select()`.
- Create new variables with functions of existing variables with `mutate()`.
- Collapse many values down to a single summary with `summarise()`.

2 / 11

## `filter()`, `arrange()` and `select()`

```
a1 <- select(airbnb, neighbourhood, price, host_id, beds, bathrooms)
```

```
a2 <- filter(a1, neighbourhood == "Richmond"  
            | neighbourhood == "Woodlawn"  
            | neighbourhood == "Downtown")
```

```
a3 <- arrange(a2, price, neighbourhood)
```

2 / 12

## `mutate()` and `transmute()`

Add new variables:

```
mutate(a3,  
       price_per_bed = price / beds,  
       price_per_bath = price / bathrooms)
```

Or, make an entirely new data frame:

```
transmute(airbnb,  
          price = price,  
          price_per_bed = price / beds,  
          price_per_bath = price / bathrooms)
```

2 / 13

## group\_by() and summarize()

group\_by() aggregates data by category, e.g.:

```
by_hood <- group_by(a3, neighbourhood)
```

Now, you can calculate *summaries* of other variables *within* each group, e.g.:

```
summarise(by_hood, price = mean(price, na.rm = TRUE))
```

2 / 14

## Your turn

1. Make a data frame only including rooms in the top ten neighbourhoods. Then, using only these neighbourhoods...
2. Find the mean price, cleaning\_fee, and ratio of cleaning fee to price, by neighbourhood.
3. Edit your code in (2) to add variables for the 25% and 75% quantile of price (use quantile( )).
4. Do as in (2) and (3) but splitting by both neighbourhood and room\_type (e.g., finding the mean price of private rooms in Woodlawn).
5. Edit your code in (1) to add a new variable giving the number of characters in the house\_rules (use nchar( )).

2 / 15

// reveal.js plugins

/